



MONASH
University

MONASH
INFORMATION
TECHNOLOGY

SQL Intermediate

FIT9132



Aggregate Functions

- COUNT, MAX, MIN, SUM, AVG
- Example:

```
SELECT max(mark)  
FROM enrolment;
```

```
SELECT min(mark)  
FROM enrolment;
```

```
SELECT avg(mark)  
FROM enrolment;
```

```
SELECT count(stu_nbr)  
FROM enrolment  
WHERE mark >= 50;
```

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q1. What will be displayed by the following SQL statement?

```
SELECT count(*), count(mark)
FROM enrolment;
```

- A. 8, 8
- B. 8, 3
- C. 3, 3
- D. 3, 8

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q2. What will be displayed by the following SQL statement?

```
SELECT count(*), count(stu_nbr), count(distinct stu_nbr)
FROM enrolment;
```

- A. 8, 8, 4
- B. 8, 8, 8
- C. 8, 4, 8
- D. 8, 4, 4


	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q3. We want to calculate the average mark of the 8 rows in the above table. What SQL statement should we use?

Note: We want to calculate $(78+35+65)/8=22.25$

- A. SELECT avg(mark) FROM enrolment;
- B. SELECT sum(mark)/count(mark) FROM enrolment;
- C. SELECT sum(mark)/count(*) FROM enrolment;
- D. SELECT avg(NVL(mark,0)) FROM enrolment;
- E. None of the above.
- F. More than one option is correct.

Anatomy of an SQL Statement - Revisited



GROUP BY

- If a GROUP BY clause is used with aggregate function, the DBMS will apply the aggregate function to the different groups defined in the clause rather than all rows.

```
SELECT avg(mark)  
FROM enrolment;
```

```
SELECT unit_code, avg(mark)  
FROM enrolment  
GROUP BY unit_code;
```

What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT avg(mark)
FROM enrolmentA;
```

```
SELECT unit_code, avg(mark)
FROM enrolmentA
GROUP BY unit_code;
```

```
SELECT unit_code, avg(mark), count(*)
FROM enrolmentA
GROUP BY unit_code;
```


What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code, year;
```

HAVING clause

- It is used to put a condition or conditions on the groups defined by GROUP BY clause.

```
SELECT unit_code, count(*)  
FROM enrolment  
GROUP BY unit_code  
HAVING count(*) > 2;
```

What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING count(*) > 2;
```

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING avg(mark) > 55;
```

HAVING and WHERE clauses

```
SELECT unit_code, count(*)  
FROM enrolment  
WHERE mark IS NULL  
GROUP BY unit_code  
HAVING count(*) > 1;
```

- The WHERE clause is applied to ALL rows in the table.
- The HAVING clause is applied to the groups defined by the GROUP BY clause.
- The order of operations performed is FROM, WHERE, GROUP BY, HAVING and then ORDER BY.
- On the above example, the logic of the process will be:
 - All rows where mark is NULL are retrieved. (due to the WHERE clause)
 - The retrieved rows then are grouped into different unit_code.
 - If the number of rows in a group is greater than 1, the unit_code and the total is displayed. (due to the HAVING clause)

What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
WHERE year = 2015  
GROUP BY unit_code  
HAVING avg(mark) > 30  
ORDER BY avg(mark) DESC;
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

Q4. What is the output for:

```
SELECT unit_code, avg(mark), studid
FROM enrolmentA
GROUP BY unit_code
HAVING avg(mark) > 55;
```

- A. FIT2094, 50, 111
- B. FIT2004, 60, 111
- C. FIT2004, 60, 111, 222, 333
- D. FIT2004, 100, 111
- E. Will print three rows
- F. Error

```
SELECT stu_lname, stu_fname, avg(mark)
FROM enrolment e JOIN student s
      ON s.stu_nbr = e.stu_nbr
GROUP BY s.stu_nbr;
```

The above SQL generates error message "ORA-00979: not a GROUP BY expression
00979. 00000 - " not a GROUP BY expression"

Why and how to fix this?

- Why? Because the grouping is based on the stu_nbr, whereas the display is based on stu_lname and stu_fname. The two groups may not have the same members.
- How to fix this?
 - Include the stu_lname,stu_fname as part of the GROUP BY condition.
- Attributes that are used in the SELECT, HAVING and ORDER BY must be included in the GROUP BY clause.

Subqueries


- Query within a query.

"Find all students whose mark is higher than the average mark of all enrolled students"


```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg (mark)  
               FROM enrolment );
```


Types of Subqueries

Single-value



Multiple-row subquery (a list of values – many rows, one column)



Multiple-column subquery (many rows, many columns)



Q5. What will be returned by the *inner query*?

```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg(mark)  
              FROM enrolment  
              GROUP BY unit_code);
```

- A. A value (a single column, single row).
- B. A list of values.
- C. Multiple columns, multiple rows.
- D. None of the above.

Q6. What will be returned by the *inner query*?

```
SELECT unit_code, stu_lname, stu_fname, mark
FROM enrolment e join student s
    on e.stu_nbr = s.stu_nbr
WHERE (unit_code, mark) IN (SELECT unit_code, max(mark)
    FROM enrolment
    GROUP BY unit_code);
```

- A. A value (a single column, single row).
- B. A list of values.
- C. Multiple columns, multiple rows.
- D. None of the above.

Comparison Operators for Subquery

- Operator for single value comparison.
=, <, >
- Operator for multiple rows or a list comparison.
 - equality
 - IN
 - inequality
 - ALL, ANY combined with <, >

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

Q7. Which row(s) in ENROL2 table will be retrieved by the following SQL statement?

```
SELECT * FROM enrol2
WHERE mark IN (SELECT max(mark)
               FROM enrol2
               GROUP BY unit_code);
```

- A. 1, 2, 7
- B. 7
- C. 2,3,7

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

Q8. Which row/s in ENROL2 will be retrieved by the following SQL statement?

```
SELECT * FROM enrol2
WHERE mark > ANY (SELECT avg(mark)
                    FROM enrol2
                    GROUP BY unit_code);
```

- A. 1,2,3,6,7
- B. 2,3,7
- C. 3,7
- D. No rows will be returned

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

Q9. Which row/s in ENROL2 will be retrieved by the following SQL statement?

```
SELECT * FROM enrol2
WHERE mark > ALL (SELECT avg(mark)
                    FROM enrol2
                    GROUP BY unit_code);
```

- A. 1,2,3,6,7
- B. 2,3,7
- C. 3,7
- D. No rows will be returned

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

Q10. Find all students whose mark in any enrolled unit is lower than 'Wendy Wheat's lowest mark of all units she is enrolled.

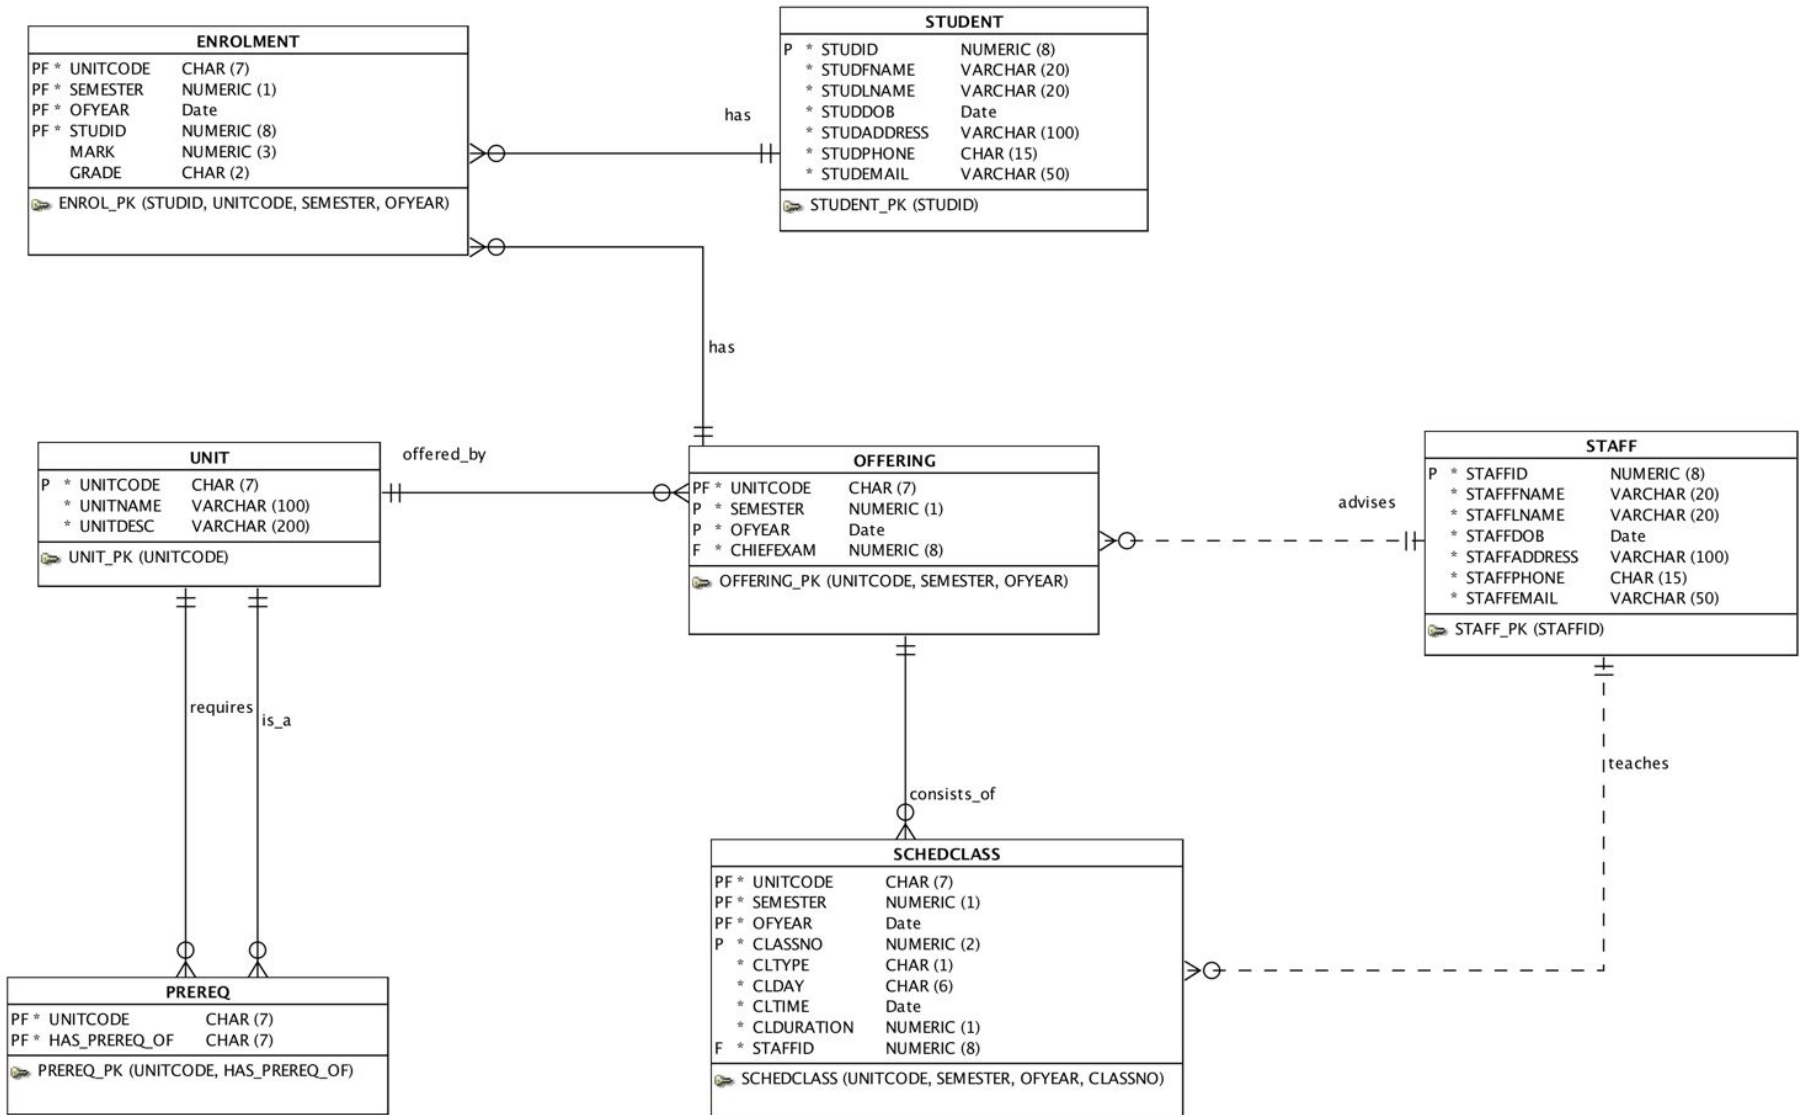
What would be a possible inner query statement for the above query?

- A. `SELECT min(mark)
FROM enrol2
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`
- B. `SELECT min(mark)
FROM enrol2 e JOIN student s on e.studid = s.studid
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`
- C. `SELECT min(mark) FROM enrol2;`
- D. `SELECT mark
FROM enrol2 e JOIN student s on e.studid = s.studid
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`

Summary

- Aggregate Functions
 - count, min, max, avg, sum
- GROUP BY and HAVING clauses.
- Subquery
 - Inner vs outer query
 - comparison operators (IN, ANY, ALL)

Practice



- Q1. Display the number of enrolments with a mark assigned, and the average mark for all enrolments
- Q2. Select the highest mark ever in any unit
- Q3. Select the highest mark ever for each unit (show the unit code only)
- Q4. For each student (show the id only), select the highest mark he/she ever received
- Q5. For each unit, print unit code, unit name and the highest mark ever in that unit. Print the results in descending order of highest marks.
- Q6. For each offering of a unit with marks show the unit code, unitname, offering details and average mark
- Q7. For each student that is enrolled in at least 3 different units, print his/her name and average mark. Also, display the number of units he/she is enrolled in.
- Q8. For each unit, count the total number of HDs
- Q9. For each unit, print the total number of HDs, Ds, and Cs. The output should contain three columns named unitcode, grade_type, num where grade_type is HD, D or C and num is the number of students that obtained the grade.
- Q10. For each unit, print the student ids of the students who obtained maximum marks in that unit